Spring Batch 2.0 Overview

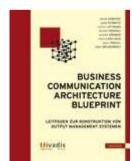


Basel · Baden Bern · Brugg · Lausanne Zurich Düsseldorf · Frankfurt/M. · Freiburg i. Br. Hamburg · Munich · Stuttgart · Vienna

Introduction

- Guido Schmutz
 - Working for Trivadis for more than 12 years
 - Co-Author of different books
 - Consultant, Trainer, Software Architect for Java, Oracle, SOA and EDA
 - Member of Trivadis Architecture Board
 - Trivadis Technology Manager
 - More than 20 years of software development experience
 - Contact: guido.schmutz@trivadis.com









Agenda



- Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Miscellaneous
- Summary

Spring Batch Introduction

- Spring Batch is the first java based framework for batch processing
 - a lightweight, comprehensive batch framework
 - builds upon the productivity, POJO-based development approach, known from the Spring Framework
 - current GA release is 1.1.4.RELEASE
 - Spring Batch 2.0 will be released in the next couple of months
 - Presentation is based on 2.0.0-RC1

Batch Processing

- What is a Batch Application?
 - Batch applications need to process high volume business critical transactional data
 - A typical batch program generally
 - 1. reads a large number of records from a database, file, or queue
 - 2. processes the data in some fashion, and
 - 3. then writes back data in a modified form

Item Oriented Processing

ItemProcessor ItemReader Step **ItemWriter** execute() read() item process(item) item read() item process(item) item write(items) ExitStatus

Usage Scenarios

- Commit batch process periodically
- Concurrent batch processing: parallel processing of a job
- Staged, enterprise message-driven processing
- Massively parallel batch processing
- Manual or scheduled restart after failure
- Sequential processing of dependent steps (with extensions to workflow-driven batches)
- Partial processing: skip records (e.g. on rollback)
- Whole-batch transaction: for cases with a small batch size or existing stored procedures/scripts

Spring Batch: Layered Architecture

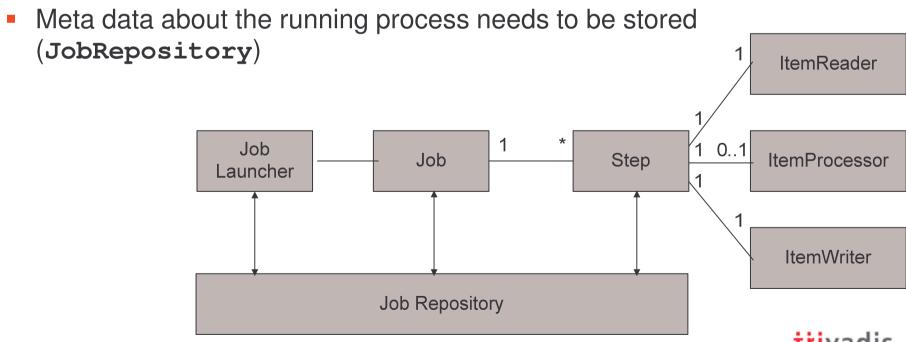
application code changes NOT REQUIRED to use different deployment approaches same application code can be used at all levels Multiple Clustered Application JVMs Multiple Workload JVMs **Batch Container** Single JVM Multi-Threaded Infrastructure Single JVM Single Process

Agenda

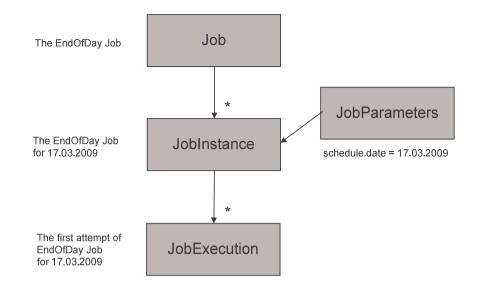


- Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Miscellaneous
- Summary

- A job has one to many steps
- A step has exactly one ItemReader, ItemWriter and optionally an ItemProcessor
- A job needs to be launched (JobLauncher)



- Job
 - encapsulates an entire batch process
- Job Instance
 - refers to the concept of a logical job run
 - job running once at end of day, will have one logical JobInstance per day
 - each JobInstance can have multiple executions
- Job Execution
 - refers to the technical concept of a single attempt to run a Job
 - An execution may end in failure or success, but the JobInstance
 - will not be considered complete unless the execution completes successfully
- Job Parameters
 - is a set of parameters used to start a batch job
 - JobInstance = Job + JobParameters



Joblnstance

*
JobExecution

*
StepExecution

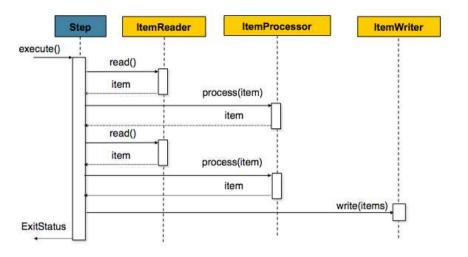
Step

- a domain object that encapsulates an independent, sequential phase of a batch job
- can be as simple or complex as the developer desires

Step Execution

- represents a single attempt to execute a Step
- A new StepExecution will be created each time a Step is run,
 similar to JobExecution
- A StepExecution will only be created when its Step is actually started

- Item Reader
 - an abstraction that represents the retrieval of input for a Step
 - one item at a time
 - When it has exhausted the items it can provide, it will indicate this by returning null
 - Various implementation available out-of-the-box
- Item Writer
 - an abstraction that represents the output of a Step
 - Chunk-oriented processing
 - Generally, an item writer has no knowledge of the input it will receive next
 - Various implementation available out-of-the-box
- Item Processor
 - an abstraction that represents the business processing of an item
 - provides access to transform or apply other business processing
 - returning null indicates that the item should not be written out



ItemReader

ItemWriter

```
public interface ItemWriter<T> {
    void write(List<? extends T> items) throws Exception;
}
```

ItemProcessor

```
public interface ItemProcessor<I, 0> {
   O process(I item) throws Exception;
}
```

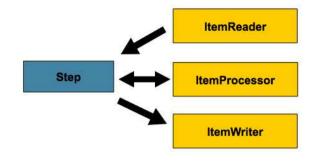
- Job Repository
 - the persistence mechanism for all of the Stereotypes
 - provides CRUD operations for JobLauncher, Job, and Step implementations
- Job Launcher
 - represents a simple interface for launching a Job with a given set of JobParameters

Agenda



- Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Miscellaneous
- Summary

Configuring and Running a Job



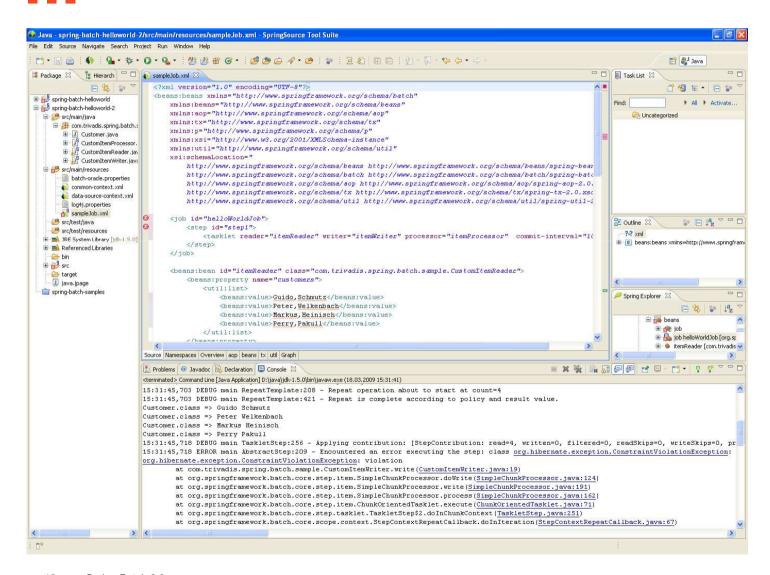
- Configuring a Job and its steps
 - There are multiple implementations of the Job interface, however, the namespace abstracts away the differences in configuration
 - It has only three required dependencies: a name, JobRepository, and a list of Steps

Configuring and Running a Job

- Configuring a Job Repository
 - used for basic CRUD operations of the various persisted domain
 objects such as JobExecution and StepExecution
 - batch namespace abstracts away many of the implementation details

Configuring a Job Launcher

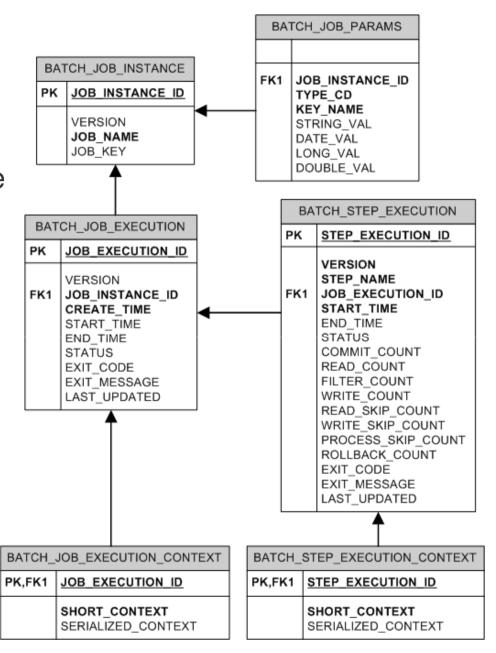
Demo





Meta-Data Schema

 The Spring Batch Meta-Data tables very closely match the Domain objects that represent them in Java

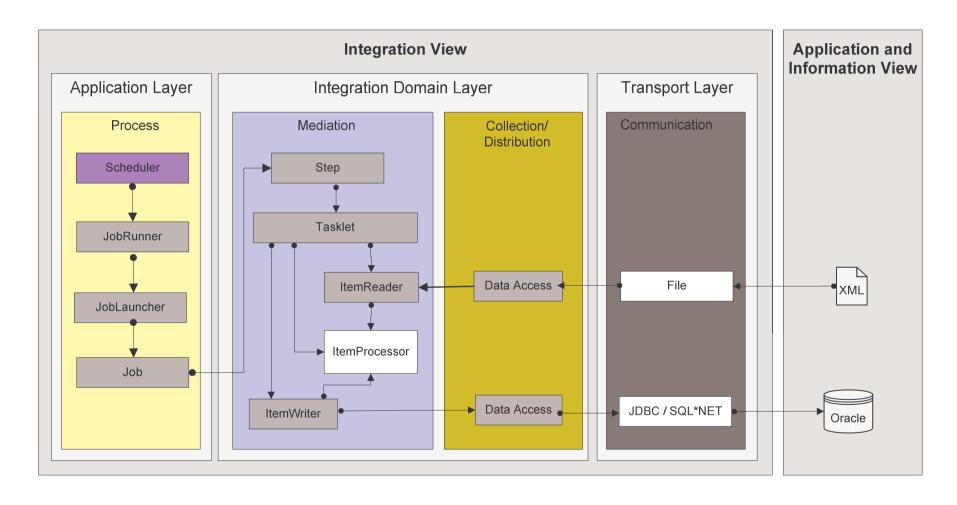


Agenda



- Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Miscellaneous
- Summary

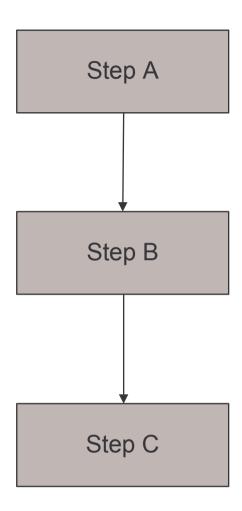
Spring Batch in Trivadis Integration Architecture Blueprint



Sequential Flow

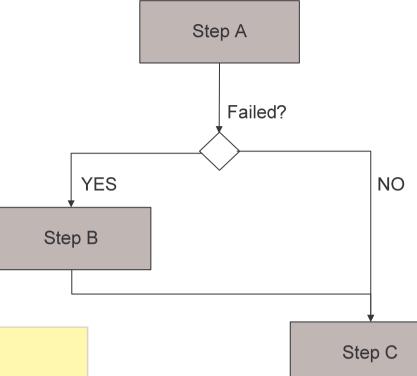
- The simplest flow scenario is a job where all of the steps execute sequentially
- This can be achieved using the 'next' attribute of the step element

```
<job id="job">
    <step id="stepA" next="stepB" />
        <step id="stepB" next="stepC"/>
        <step id="stepC"/>
        </job>
```



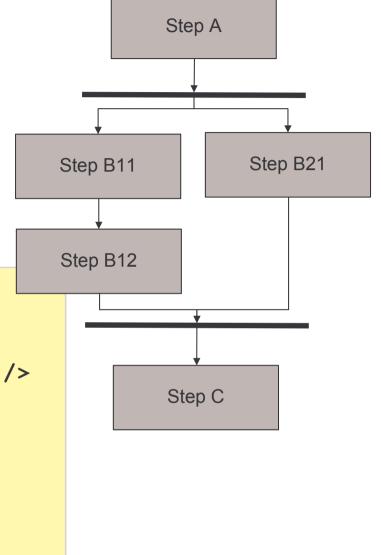
Conditional Flow

 In order to handle more complex scenarios, Spring Batch allows transition elements to be defined within the step element



Split Flow

 Spring Batch also allows for a job to be configured with parallel flows using the 'split' element



Restartablility

- The launching of a Job is considered to be a 'restart' if a JobExecution already exists for the particular JobInstance.
 - Ideally, all jobs should be able to start up where they left off
 - but there are scenarios where this is not possible
- If a Job should never be restarted, but should always be run as part of a new JobInstance, then the restartable property may be set to 'false'

```
<job id="footballJob" restartable="false">
        <step id="playerload" next="gameLoad"/>
        <step id="gameLoad" next="playerSummarization"/>
        <step id="playerSummarization"/>
        </job>
```

Configuring a Step for Restart

- Setting a StartLimit
 - control the number of times a Step may be started

- Restarting a completed step
 - In a restartable job, one or more steps should always be run, regardless of whether or not they were successful the first time

Configuring Skip Logic

 there are scenarios where errors encountered should not result in Step failure, but should be skipped instead

Configuring Fatal Exceptions

 it may be easier to identify which exceptions should cause failure and skip everything else

Intercepting Step Execution

 You might need to perform some functionality at certain events during the execution of a Step

- can be accomplished with one of many Step scoped listeners, like
 - StepExecutionListener
 - ChunkListener
 - ItemReadListener
 - ItemProcessListener
 - ItemWriteListener
 - SkipListener



Available Item Readers

Item Reader	Description		
AbstractItemCountingItemStreamItemReader	Abstract base class that provides basic restart capabilities by counting the number of items returned from an ItemReader.		
ListItemReader	Provides the items from a list, one at a time		
ItemReaderAdapter	Adapts any class to the ItemReader interface.		
AggregateItemReader	An ItemReader that delivers a list as its item, storing up objects from the injected ItemReader until they are ready to be packed out as a collection. This ItemReader should mark the beginning and end of records with the constant values in FieldSetMapper AggregateItemReader#BEGIN_RECORD and AggregateItemReader#END_RECORD		
FlatFileItemReader	Reads from a flat file. Includes ItemStream and Skippable functionality. See section on Read from a File		
StaxEventItemReader	Reads via StAX. See HOWTO - Read from a File		
JdbcCursorItemReader	Reads from a database cursor via JDBC. See HOWTO - Read from a Database		
HibernateCursorItemReader	Reads from a cursor based on an HQL query. See section on Reading from a Database		
IbatisPagingItemReader	Reads via iBATIS based on a query. Pages through the rows so that large datasets can be read without running out of memory. See HOWTO - Read from a Database		
JmsItemReader	Given a Spring JmsOperations object and a JMS Destination or destination name to send errors, provides items received through the injected JmsOperations receive() method		
JpaPagingItemReader	Given a JPQL statement, pages through the rows, such that large datasets can be read without running out of memory		
JdbcPagingItemReader	Given a SQL statement, pages through the rows, such that large datasets can be read without running out of memory		



Available Item Writers

Item Writer	Description
AbstractItemStreamItemWriter	Abstract base class that combines the ItemStream and ItemWriter interfaces.
CompositeItemWriter	Passes an item to the process method of each in an injected List of ItemWriter objects
ItemWriterAdapter	Adapts any class to the ItemWriter interface.
PropertyExtractingDelegatingItemWriter	Extends AbstractMethodInvokingDelegator creating arguments on the fly. Arguments are created by retrieving the values from the fields in the item to be processed (via a SpringBeanWrapper) based on an injected array of field name
FlatFileItemWriter	Writes to a flat file. Includes ItemStream and Skippable functionality. See section on Writing to a File
HibernateItemWriter	This item writer is hibernate session aware and handles some transaction-related work that a non-"hibernate aware" item writer would not need to know about and then delegates to another item writer to do the actual writing.
JdbcBatchItemWriter	Uses batching freatures from a PreparedStatement, if available, and can take rudimentary steps to locate a failure during a flush.
JpaItemWriter	This item writer is JPA EntityManager aware and handles some transaction-related work that a non-"jpa aware" ItemWriter would not need to know about and then delegates to another writer to do the actual writing.
StaxEventItemWriter	Uses an ObjectToXmlSerializer implementation to convert each item to XML and then writes it to an XML file using StAX.



Agenda



- Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Miscellaneous
- Summary

Summary

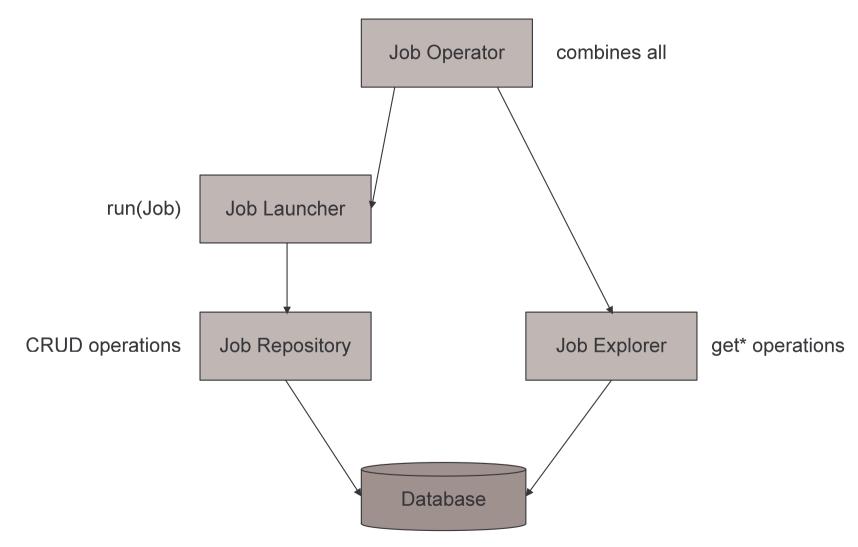
- Lack of a standard enterprise batch architecture is resulting in higher costs associated with the quality and delivery of solutions.
- Spring Batch provides a highly scalable, easy-to-use, customizable, industry-accepted batch framework collaboratively developed by Accenture and SpringSource
- Spring patterns and practices have been leveraged allowing developers to focus on business logic, while enterprise architects can customize and extend architecture concerns



Thank you!



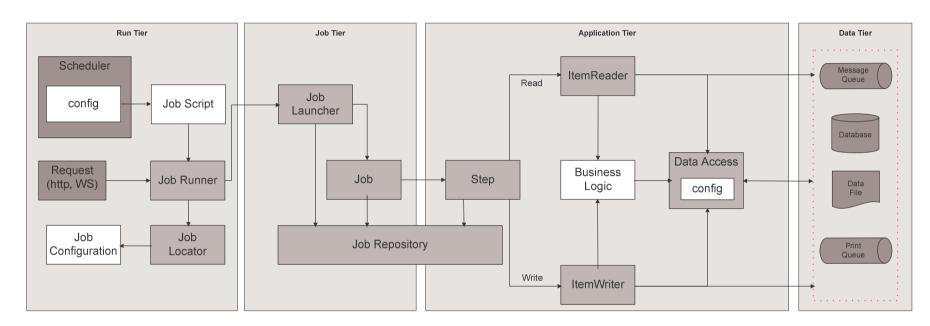
Basel · Baden Bern · Brugg · Lausanne Zurich Düsseldorf · Frankfurt/M. · Freiburg i. Br. Hamburg · Munich · Stuttgart · Vienna



JobRuntimeInformation BatchContainerLauncher uses to identify and manage jobs starts & stops uses to construct jobs BatchContainer Application Developer recipe for executes Job JobExecutor configures JobConfiguration locates * executes Step StepExecutor StepConfiguration implements stored in uses Module??? RepeatTemplate Database ExceptionHandler uses CompletionPolicy other infrastructure dependencies



Spring Batch Launch Environment





Spring Batch Launch Environment

Run Tier

- concerned with the scheduling and launching of the application
- a vendor product is typically used in this tier to allow
 - time-based and interdependent scheduling of batch jobs
 - providing parallel processing capabilities

Job Tier

- responsible for the overall execution of a batch job
- sequentially executes batch steps, ensuring that all steps are in the correct state and all appropriate policies are enforced

Application Tier

- contains components required to execute the program
- contains specific modules that address the required batch functionality and enforces policies around a module execution (e.g., commit intervals, capture of statistics, etc.)

Data Tier

 provides the integration with the physical data sources that might include databases, files, or queues.