



Job Scheduling with Spring

Most application logic happens in response to some form of user action, such as a button click or a form submission. However, in many applications certain processes must be invoked *without* user interaction, usually at a given interval. For example, you might have a process that cleans out temporary files once an hour or a process that creates a data export from a database and sends it to an external system once a day at midnight. Most nontrivial applications require some kind of scheduling support—if not directly related to business logic of the application then to support system housekeeping.

When you are building scheduled tasks for your application, creating a task that runs once an hour or once a day is fairly simple. But what about a task that runs at 3:00 p.m. every Monday, Wednesday, and Friday? This is a little more difficult to code, and it makes sense to use prebuilt scheduling solutions rather than attempt to create your own scheduling framework.

When talking about scheduling from a programming perspective, we tend to talk about three distinct concepts. A job is a unit of work that needs to be scheduled to run at a specific interval. A trigger is a condition, perhaps a fixed interval or a given piece of data, that causes a job to run. A schedule is a collection of triggers that govern the complete timeline of a job. Typically, you encapsulate a job by implementing an interface or extending a given base class. You define your triggers in whatever terms your scheduling framework supports. Some frameworks may support only basic interval-based triggers, whereas others, such as Quartz, provide much more flexible trigger schemes. In general, a job has only a single trigger in its schedule, and the terms “schedule” and “trigger” are often used interchangeably.

Scheduling support in Spring comes in two distinct forms: JDK Timer-based and Quartz-based. The JDK Timer-based approach provides scheduling capabilities on any version 1.3 or later JVM, and it does not need external dependencies beyond Spring. Timer-based scheduling is quite primitive and provides limited flexibility when defining job schedules. However, Timer support is included with Java and requires no external libraries, which might be beneficial if you are restricted by application size or corporate policy. Quartz-based scheduling is much more flexible and allows triggers to be defined in a much more real-world way, such as the earlier example of 3:00 p.m. every Monday, Wednesday, and Friday.

In this chapter, we explore both of the scheduling solutions included with Spring. In particular, this chapter discusses three core topics: scheduling with the JDK Timer, Quartz-based scheduling, and job scheduling considerations.

We'll start our discussion by exploring Spring's support for JDK Timer-based scheduling. This section introduces the different trigger types available with Timer-based scheduling and looks at how you can schedule any arbitrary logic without needing to create additional Java code.

After that, we'll look at the comprehensive Quartz scheduling engine and how it is integrated into Spring. In particular, we examine Quartz support for cron expressions allowing highly complex schedules to be configured using a concise format. As with the JDK Timer, you'll see how to schedule any logic without needing to encapsulate it.

Finally, we'll discuss the various points to consider when choosing a scheduling implementation and patterns to use when creating logic for scheduled execution.

Scheduling Jobs Using JDK Timer

The most basic scheduling support with Spring is based on the JDK `java.util.Timer` class. When scheduling using `Timer`, you are limited to simple interval-based trigger definitions, which makes `Timer`-based scheduling suitable only for jobs that you need to execute just once at some given future time or that you need to execute at some fixed frequency.

Trigger Types with Timer

`Timer`-based scheduling offers you three types of triggers:

- **One-off:** When you use a one-off trigger, job execution is scheduled for some given point in the future, defined as the number of milliseconds from a given date. After the job executes, it is not rescheduled for further execution. We have found that one-off triggers are great for scheduling jobs that need to be done once that you might forget to do yourself. For instance, if a web application has scheduled maintenance coming up in a week, we can schedule a task to switch on the In Maintenance page when the maintenance is due to begin.
- **Repeating and fixed-delay:** When you use a fixed-delay trigger, you schedule the first execution of the job just like for a one-off trigger, but after that, the job is rescheduled to execute after a given interval. When you are using fixed-delay, the interval is relative to the actual execution time of the previous job. This means that the interval between successive executions is always approximately the same, even if executions occur “late” when compared to the original schedule. With this type of trigger, the interval you specify is the actual interval between subsequent executions. Use this approach when the interval between executions must be kept as constant as possible.
- **Repeating and fixed-rate:** Fixed-rate triggers function in a similar way to fixed-delay triggers, but the next execution time is always calculated based on the initial scheduled execution time. This means that if a single execution is delayed, subsequent executions are not delayed as a result. With this type of trigger, the interval you specify is not necessarily the actual interval between subsequent executions. Use this approach when the actual execution time is important, rather than the interval between executions.

You may find it difficult to visualize the differences between fixed-delay and fixed-rate triggers, and unfortunately, creating an example that causes enough of a delay in execution to fully clarify the differences reliably is difficult. That said, here is a simple example that should help highlight the differences.

Consider a task that starts executing at 1:00 p.m. and has a specified interval of 30 minutes. The task runs fine until 4:30 p.m., when the system experiences a heavy load and a particularly nasty garbage collection; these cause the actual execution time to be a minute late—4:31 p.m. Now, with fixed-delay scheduling it is the *actual interval* that is important, that is to say, we want 30 minutes between each actual execution, so the next execution is scheduled for 5:01 p.m. rather than 5:00 p.m. With fixed-rate scheduling, the interval defines the *intended* interval—that is to say, we intend the job to execute every 30 minutes based on the start time, not on the time of the last job—so the job is scheduled for execution at 5:00 p.m.

Both of these trigger types are useful in different ways. In general, you use fixed-delay triggers for situations where you want the time between each execution to be as regular as possible or when you want to avoid the possibility of two executions happening too close together, which can happen

with fixed-rate execution if a particular execution is delayed long enough. You use fixed-rate triggers for real-time-sensitive operations such as those that must execute every hour on the hour.

Creating a Simple Job

To create a job to use with the `Timer` class, you simply extend the `TimerTask` class and implement the `run()` method to execute your job's logic. Listing 12-1 shows a simple `TimerTask` implementation that writes “Hello, World” to `stdout`.

Listing 12-1. *Creating a Basic `TimerTask`*

```
package com.apress.prospring2.ch12.timer;

import java.util.TimerTask;

public class HelloWorldTask extends TimerTask {

    public void run() {
        System.out.println("Hello World!");
    }
}
```

Here, you can see that in the `run()` method, we simply write the “Hello, World” message to `stdout`. Each time a job is executed, `Timer` invokes the `TimerTask`'s `run()` method. The simplest possible trigger we can create for this job is a one-off trigger to start the job in 1 second; Listing 12-2 shows this.

Listing 12-2. *Using a One-Off Trigger with the `HelloWorldTask`*

```
package com.apress.prospring2.ch12.timer;

import java.util.Timer;

public class OneOffScheduling {

    public static void main(String[] args) {
        Timer t = new Timer();
        t.schedule(new HelloWorldTask(), 1000);
    }
}
```

To schedule a job using a given trigger when you are using the JDK `Timer` class, you must first create an instance of the `Timer` class and then create the trigger using one of the `schedule()` or `scheduleAtFixedRate()` methods. In Listing 12-2, we used the `schedule()` method to schedule an instance of `HelloWorldTask` to run after a delay of 1,000 milliseconds. If you run this example, after the initial delay of 1 second, you get the following message:

```
Hello World!
```

This kind of one-off trigger is fairly useless—how often are you going to need to schedule a one-off task to run an arbitrary period of time after application start-up? For this reason, you can also specify an absolute date when you create a one-off trigger. So if we want to create a job to remind us seven days before an important birthday, we can replace our call to `Timer.schedule()` with something like this:

```
Calendar cal = Calendar.getInstance();
cal.set(2008, Calendar.NOVEMBER, 30);
t.schedule(new HelloWorldTask(), cal.getTime());
```

In this example, you can see that we created an instance of `Calendar` for the date November 30, 2008. Then, using the `Calendar` instance, we scheduled the `HelloWorldTask` to run. This is clearly more useful than the first example, because no matter what time the application starts, the job is always scheduled to run at the same time. The only drawback with this approach is that we will not be reminded about the birthday in 2009 or 2010 unless we explicitly add more triggers. By using a repeating trigger, we can get around this.

Both types of repeating trigger, fixed-delay and fixed-rate, are configured in the same way: you specify a starting point, using either a number of milliseconds relative to the call to `schedule()` or an absolute date, and you specify an interval in milliseconds to control when subsequent executions occur. Remember that “interval” is interpreted differently depending on whether you are using a fixed-delay or fixed-rate trigger.

We can schedule the `HelloWorldTask` job to run every 3 seconds with a 1-second delay using the code shown in Listing 12-3.

Listing 12-3. *Scheduling a Repeating Task*

```
package com.apress.prospring2.ch12.timer;

import java.util.Timer;

public class FixedDelayScheduling {

    public static void main(String[] args) throws Exception{
        Timer t = new Timer();
        t.schedule(new HelloWorldTask(), 1000, 3000);
    }
}
```

If you run this application, you will see the first “Hello, World” message displayed after about 1 second, followed by further “Hello, World” messages every 3 seconds. To schedule this job using a fixed-rate trigger, simply replace the call to `Timer.schedule()` with a call to `Timer.scheduleAtFixedRate()`, as shown in Listing 12-4.

Listing 12-4. *Scheduling a Job Using a Fixed-Rate Trigger*

```
package com.apress.prospring2.ch12.timer;

import java.util.Timer;

public class FixedRateScheduling {

    public static void main(String[] args) throws Exception {
        Timer t = new Timer();
        t.scheduleAtFixedRate(new HelloWorldTask(), 1000, 1000);
    }
}
```

As with the one-off trigger, you can start both fixed-delay and fixed-rate triggers using an absolute date. Using this approach, we can create a trigger for our birthday reminder example that runs on a given date and then repeats each year. This is shown in Listing 12-5.

Listing 12-5. *Scheduling Birthday Reminders*

```
package com.apress.prospring2.ch12.timer;

import java.util.Calendar;
```

```
import java.util.Timer;

public class SimpleBirthdayReminderScheduling {

    private static final long MILLIS_IN_YEAR = 1000 * 60 * 60 * 24 * 365;

    public static void main(String[] args) {
        Timer t = new Timer();

        Calendar cal = Calendar.getInstance();
        cal.set(2008, Calendar.NOVEMBER, 30);
        t.schedule(new HelloWorldTask(), cal.getTime());

        t.scheduleAtFixedRate(new HelloWorldTask(), cal.getTime(),
                               MILLIS_IN_YEAR);
    }
}
```

In this example, you can see that we calculate the number of milliseconds in a year, and using a `Calendar` instance, we define a starting point of November 30 and define the interval to be one year. Now, every year on November 30, provided that this application is running and conveniently ignoring the existence of leap years, the “Hello, World” message is written to `stdout`. Clearly, this is not a fully functional example, because there is no real notification mechanism, and each time we want to add a new birthday reminder, we need to change the code. In the next section, we create a more robust birthday reminder application using Spring’s JDK Timer support classes.

Spring Support for JDK Timer Scheduling

As you saw in the previous section, you can easily create and schedule jobs using the JDK `Timer` and `TimerTask` classes. That said, the approach we took in the previous examples has some problems. First, we created the `TimerTask` instances within the application rather than using Spring. For the `HelloWorldTask`, this is acceptable, because we did not need to configure the job at all. However, many jobs require some configuration data, so we should manage these using Spring to allow for easy configuration. Second, the trigger information is hard-coded into the application, meaning that any changes to the time a job is triggered require a change to the application code and a recompilation. Finally, scheduling new jobs or removing a job requires changes to the application code; ideally, we should be able to configure this externally. By using Spring’s Timer support classes, we can externalize all job and trigger configuration as well as hand over control of `Timer` creation to Spring, thus allowing jobs and their triggers to be defined externally.

The core of Spring’s Timer support comes in the form of the `ScheduledTimerTask` and `TimerFactoryBean` classes. The `ScheduledTimerTask` class acts as a wrapper around your `TimerTask` implementations and allows you to define trigger information for the job. Using the `TimerFactoryBean`, you can have Spring automatically create `Timer` instances for a given list of `ScheduledTimerTask` beans using the trigger configuration data when creating the trigger.

Using `ScheduledTimerTask` and `TimerFactoryBean`

Before we dive in and look at our new and improved birthday reminder application, we should first look at the basics of how `ScheduledTimerTask` and `TimerFactoryBean` work. For each scheduled job you want to create, you need to configure an instance of the job class and an instance of `ScheduledTimerTask` containing the trigger details. You can share the same `TimerTask` instance across many `ScheduledTimerTask` instances if you want to create many triggers for the same job. Once you have these components configured, simply configure a `TimerFactoryBean` and specify

the list of `ScheduledTimerTask` beans. Spring then creates an instance of `Timer` and schedules all the jobs defined by the `ScheduledTimerTask` beans using that `Timer` class.

This might sound complex at first, but in reality, it is not. Listing 12-6 shows a simple configuration for scheduling the `HelloWorldTask` to run every 3 seconds with a delay of 1 second before the first execution.

Listing 12-6. *Configuring Job Scheduling Using `TimerFactoryBean`*

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns=http://www.springframework.org/schema/beans
       xmlns:xsi=http://www.w3.org/2001/XMLSchema-instance
       xsi:schemaLocation="http://www.springframework.org/schema/beans
       http://www.springframework.org/schema/beans/spring-beans-2.0.xsd">

    <bean id="job" class="com.apress.prospring2.ch12.timer>HelloWorldTask"/>

    <bean id="timerTask"
        class="org.springframework.scheduling.timer.ScheduledTimerTask">
        <property name="delay" value="1000" />
        <property name="period" value="3000" />
        <property name="timerTask" ref="job" />
    </bean>

    <bean id="timerFactory"
        class="org.springframework.scheduling.timer.TimerFactoryBean">
        <property name="scheduledTimerTasks">
            <list>
                <ref local="timerTask"/>
            </list>
        </property>
    </bean>
</beans>
```

Here, you can see that we have configured a bean, `job`, of type `HelloWorldTask`, and using this bean, we have configured a bean of type `ScheduledTimerTask`, setting the delay to 1,000 milliseconds and the period to 3,000 milliseconds. The final part of the configuration is the `timerFactory` bean, which is passed a list of beans of type `ScheduledTimerTask`. In this case, we have only one task to schedule, represented by the `timerTask` bean. Be aware that when specifying trigger information using `ScheduledTimerTask`, you can supply a delay only in milliseconds, not an initial date for start-up. We'll show you a way around this when we build the birthday reminder application in the next section.

With all of the scheduling and job definition information contained in the configuration, our sample application has very little to do. In fact, all we need to do is load the `ApplicationContext`, and Spring creates the `Timer` class and schedules the `HelloWorldTask` as per the configuration file. This code is shown in Listing 12-7.

Listing 12-7. *The `TimerFactoryBeanExample` Class*

```
package com.apress.prospring2.ch12.timer;

import org.springframework.context.ApplicationContext;
import org.springframework.context.support.FileSystemXmlApplicationContext;

public class TimerFactoryBeanExample {
```

```

    public static void main(String[] args) throws Exception {
        ApplicationContext ctx = new FileSystemXmlApplicationContext(
            "./ch12/src/conf/timer-context.xml");
        System.in.read();
    }
}

```

If you run this application, you will see that the message “Hello, World” is written to stdout every 3 seconds after an initial delay of 1 second. As you can see from this example, it is very simple to configure job scheduling external to your application's code. Using this approach, it is much simpler to make changes to a job's schedules or to add new scheduled jobs and remove existing ones.

A More Comprehensive Birthday Reminder Application

In this section, we create a more complex birthday reminder application using Spring's Timer support. With this example, we want to be able to schedule multiple reminder jobs, each with a specific configuration, to identify whose birthday the reminder indicates. We also want to be able to add and remove reminders without having to modify the application code.

To get started, we need to create a job to perform the actual reminder. Because we are going to create these jobs using Spring, we can allow all configuration data to be provided using DI. Listing 12-8 shows the `BirthdayReminderTask`.

Listing 12-8. *The BirthdayReminderTask*

```

package com.apress.prospring2.ch12.timer.bday;

import java.util.TimerTask;

public class BirthdayReminderTask extends TimerTask {

    private String who;

    public void setWho(String who) {
        this.who = who;
    }

    public void run() {
        System.out.println("Don't forget it is " + who
            + "'s birthday is 7 days");
    }
}

```

Notice here that we defined a property on the task, `who`, that allows us to specify of whose birthday we're being reminded. In a real birthday reminder application, the reminder would no doubt be sent to e-mail or some similar medium. For now, however, you'll have to be content with reminder messages written to stdout!

With this task complete, we are almost ready to move on to the configuration stage. However, as we pointed out earlier, you cannot specify the start time of a scheduled job using a date when you are using `ScheduledTimerTask`. This is problematic for our sample application, because we do not want to have to specify reminder dates as a relative offset to the start-up time of the application. Thankfully, we can overcome this problem quite easily by extending the `ScheduledTimerTask` class and overriding the `getDelay()` method used by `TimerFactoryBean` to determine what delay it should assign to a trigger. At the same time, we can also override the `getPeriod()` method to return the number of milliseconds in a year so that you do not have to add that literal into configuration files. Listing 12-9 shows the code for our custom `ScheduledTimerTask`, `BirthdayScheduledTask`.

Listing 12-9. *Customizing ScheduledTimerTask*

```

package com.apress.prospring2.ch12.timer.bday;

import java.text.DateFormat;
import java.text.ParseException;
import java.text.SimpleDateFormat;
import java.util.Calendar;
import java.util.Date;

import org.springframework.scheduling.timer.ScheduledTimerTask;

public class BirthdayScheduledTask extends ScheduledTimerTask {

    private static final long MILLIS_IN_YEAR = 1000 * 60 * 60 * 24 * 365;

    private DateFormat dateFormat = new SimpleDateFormat("yyyy-MM-dd");

    private Date startDate;

    public void setDate(String date) throws ParseException {
        startDate = dateFormat.parse(date);
    }

    public long getDelay() {
        Calendar now = Calendar.getInstance();
        Calendar then = Calendar.getInstance();
        then.setTime(startDate);

        return (then.getTimeInMillis() - now.getTimeInMillis());
    }

    public long getPeriod() {
        return MILLIS_IN_YEAR;
    }
}

```

In this example, you can see that we define a new property for the `BirthdayScheduledTask` class, `date`, which allows us to specify the start date as a date rather than a delay period. This property is of type `String`, because we use an instance of `SimpleDateFormat` configured with the pattern `yyyy-MM-dd` to parse dates such as `2008-11-30`. We override the `getPeriod()` method, which `TimerFactoryBean` uses when it configures the interval for the trigger, to return the number of milliseconds in a year. Also notice that we override `getDelay()`, and using the `Calendar` class, we calculate the number of milliseconds between the current time and the specified start date. This value is then returned as the delay. With this complete, we can now complete the configuration for our sample application, as shown in Listing 12-10.

Listing 12-10. *Configuring the Birthday Reminder Application*

```

<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xsi:schemaLocation="http://www.springframework.org/schema/beans
        http://www.springframework.org/schema/beans/spring-beans-2.0.xsd">

```



```

<bean id="mum"
    class="com.apress.prospring2.ch12.timer.bday.BirthdayScheduledTask">
    <property name="date" value="2008-11-30" />
    <property name="fixedRate" value="true" />

    <property name="timerTask">
        <bean class="com.apress.prospring2.ch12.timer.bday.BirthdayReminderTask">
            <property name="who" value="Mum">
            </bean>
        </property>
    </property>
</bean>

<bean id="timerFactory"
    class="org.springframework.scheduling.timer.TimerFactoryBean">
    <property name="scheduledTimerTasks">
        <list>
            <ref local="mum"/>
        </list>
    </property>
</bean>
</beans>

```

This code should look familiar to you by now. Notice that we used our `BirthdayScheduledTask` class in place of the `ScheduledTimerTask` class, and instead of specifying a delay and a period, we have simply specified the date. Also, we rely on the overridden `getDelay()` and `getPeriod()` methods to provide the `TimerFactoryBean` with the delay and period values. In addition, notice that we set the `fixedRate` property of the `BirthdayScheduledTask` bean to `true`. This property is inherited from `ScheduledTimerTask`; `TimerFactoryBean` uses it to decide whether or not it should create a fixed-rate or fixed-delay trigger.

Scheduling Arbitrary Jobs

When scheduling jobs, you often need to schedule the execution of logic that already exists. If this is the case, you might not want to go through the trouble of creating a `TimerTask` class just to wrap your logic. Thankfully, you don't have to. Using the `MethodInvokingTimerTaskFactoryBean`, you can schedule the execution of any method on any given bean or a static method on a specific class; you can even provide method arguments if your logic method requires them.

As an example of this, consider the `FooBean` shown in Listing 12-11.

Listing 12-11. The `FooBean` Class

```

package com.apress.prospring2.ch12.timer;

public class FooBean {

    public void someJob(String message) {
        System.out.println(message);
    }
}

```

If we want to schedule the `someJob()` method to run every 3 seconds with a given argument rather than create a `TimerTask` just to do that, we can simply use the `MethodInvokingTimerTaskFactoryBean` to create a `TimerTask` for us. The configuration for this is shown in Listing 12-12.

Listing 12-12. *Using MethodInvokingTimerTaskFactoryBean*

```

<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns=http://www.springframework.org/schema/beans
       xmlns:xsi=http://www.w3.org/2001/XMLSchema-instance
       xsi:schemaLocation="http://www.springframework.org/schema/beans
       http://www.springframework.org/schema/beans/spring-beans-2.0.xsd">

    <bean id="target" class="com.apress.prospring2.ch12.timer.FooBean"/>

    <bean id="task" class="org.springframework.scheduling.timer.➤
                                   MethodInvokingTimerTaskFactoryBean">
        <property name="targetObject" ref="target" />
        <property name="targetMethod" value="someJob" />
        <property name="arguments" value="Hello World!" />
    </bean>

    <bean id="timerTask"
          class="org.springframework.scheduling.timer.ScheduledTimerTask">
        <property name="delay" value="1000" />
        <property name="period" value="3000" />
        <property name="timerTask" ref="task" />
    </bean>

    <bean id="timerFactory"
          class="org.springframework.scheduling.timer.TimerFactoryBean">
        <property name="scheduledTimerTasks">
            <list>
                <ref local="timerTask"/>
            </list>
        </property>
    </bean>

</beans>

```

We can replace the definition of our own custom `TimerTask` bean with a definition using the `MethodInvokingTimerTaskFactoryBean`. To configure `MethodInvokingTimerTaskFactoryBean`, we specify the target of the invocation as a reference to another bean, the method to execute, and the argument to use when executing. The `TimerTask` supplied by `MethodInvokingTimerTaskFactoryBean` is used in the normal way, wrapped in a `ScheduledTimerTask`, and passed to the `TimerFactoryBean`.

Listing 12-13 shows a simple driver program to test this out.

Listing 12-13. *Testing the MethodInvokingTimerTaskFactoryBean*

```

package com.apress.prospring2.ch12.timer;

import org.springframework.context.ApplicationContext;
import org.springframework.context.support.FileSystemXmlApplicationContext;

public class MethodInvokerScheduling {

    public static void main(String[] args) throws Exception {
        ApplicationContext ctx = new FileSystemXmlApplicationContext(
            "./ch12/src/conf/timerMethodInvoker.xml");
        System.in.read();
    }
}

```

Running this example gives you the now familiar timed appearance of “Hello, World” messages on your console. Clearly, using `MethodInvokingTimerTaskFactoryBean` removes the need to create custom `TimerTask` implementations that simply wrap the execution of a business method.

JDK `Timer`-based scheduling provides support for an application’s basic scheduling needs using a simple and easy-to-understand architecture. Although the trigger system for JDK `Timer` is not extremely flexible, it does provide basic schemes that allow you to control simple scheduling. Using Spring’s support classes for `Timer`, you externalize a task scheduling configuration and make it easier to add and remove tasks from the scheduler without having to change any application code. Using `MethodInvokingTimerTaskFactoryBean`, you avoid having to create `TimerTask` implementations that do nothing more than invoke a business method, thus reducing the amount of code you need to write and maintain.

The main drawback of JDK `Timer` scheduling comes when you need to support complex triggers such as a trigger to execute a job every Monday, Wednesday, and Friday at 3:00 p.m. In the next part of this chapter, we look at the Quartz scheduling engine, which provides much more comprehensive support for scheduling and, just like `Timer`, is fully integrated into Spring.

Scheduling Jobs Using OpenSymphony Quartz

The open source Quartz project is a dedicated job scheduling engine designed to be used in both Java EE and Java SE settings. Quartz provides a huge range of features such as persistent jobs, clustering, and distributed transactions, though we do not look at the clustering or distributed transaction features in this book—you can find out more about these online at www.opensymphony.com/quartz. Spring’s Quartz integration is similar to its `Timer` integration in that it provides for declarative configuration of jobs, triggers, and schedules. In addition to this, Spring provides additional job persistence features that allow the scheduling of a Quartz job to take part in a Spring-managed transaction.

Introducing Quartz

Quartz is an extremely powerful job scheduling engine, and we cannot hope to explain everything about it in this chapter. However, we do cover the main aspects of Quartz that are related to Spring, and we discuss how you can use Quartz from a Spring application. As with our `Timer` discussion, we start by looking at Quartz separately from Spring, and then we look at integrating Quartz and Spring.

The core of Quartz is made up of two interfaces, `Job` and `Scheduler`, and two classes, `JobDetail` and `Trigger`. From their names, it should be apparent what `Job`, `Scheduler`, and `Trigger` do, but the role of the `JobDetail` class is not so clear. Unlike `Timer`-based scheduling, tasks are not executed using a single instance of your job class; instead, Quartz creates instances as it needs them. You can use the `JobDetail` class to encapsulate the job state and to pass information to a job between subsequent executions of a job. `Timer`-based scheduling has no notion of a `Trigger` class; `Trigger` logic is encapsulated by the `Timer` class itself. Quartz supports a pluggable architecture for triggers, which allows you to create your own implementations as you see fit. That said, you rarely create your own `Trigger` implementations because Quartz provides the superpowerful `CronTrigger` class out of the box, which allows you to use cron expressions (more on that shortly) to have fine-grained control over job execution.

Simple Job Scheduling

To create a job for use in Quartz, simply create a class that implements the `Job` interface. The `Job` interface defines a single method, `execute()`, from which you call your business logic. Quartz passes an instance of `JobExecutionContext` to the `execute()` method, allowing you to access context data about the current execution. We’ll look at this in more detail in the next section.

Listing 12-14 show a simple `Job` implementation that writes “Hello, World” to `stdout`.

Listing 12-14. *Creating a Simple Job*

```
package com.apress.prospring2.ch12.quartz;

import org.quartz.Job;
import org.quartz.JobExecutionContext;
import org.quartz.JobExecutionException;

public class HelloWorldJob implements Job {

    public void execute(JobExecutionContext context)
        throws JobExecutionException {
        System.out.println("Hello World!");
    }

}
```

To schedule this job to run, we first need to obtain a Scheduler instance, then create a JobDetail bean that contains information about the job, and finally create a Trigger to govern job execution. The code for this is shown in Listing 12-15.

Listing 12-15. *Scheduling Jobs in Quartz*

```
package com.apress.prospring2.ch12.quartz;

import java.util.Date;

import org.quartz.JobDetail;
import org.quartz.Scheduler;
import org.quartz.SimpleTrigger;
import org.quartz.Trigger;
import org.quartz.impl.StdSchedulerFactory;

public class HelloWorldScheduling {

    public static void main(String[] args) throws Exception {

        Scheduler scheduler = new StdSchedulerFactory().getScheduler();
        scheduler.start();

        JobDetail jobDetail = new JobDetail("helloWorldJob",
            Scheduler.DEFAULT_GROUP, HelloWorldJob.class);

        Trigger trigger = new SimpleTrigger("simpleTrigger",
            Scheduler.DEFAULT_GROUP, new Date(), null,
            SimpleTrigger.REPEAT_INDEFINITELY, 3000);

        scheduler.scheduleJob(jobDetail, trigger);
    }

}
```

This code starts by obtaining an instance of Scheduler using the StdSchedulerFactory class. We are not going to look at this class in any detail here, but you can find out more information in the Quartz tutorial, which is available on the OpenSymphony web site. For now, it is enough to know that the StdSchedulerFactory.getScheduler() class returns a Scheduler instance that is ready to run. In Quartz, a Scheduler can be started, stopped, and paused. If a Scheduler has not been started or is paused, no triggers fire, so we start the Scheduler using the start() method.

Next, we create the `JobDetail` instance of the job we are scheduling, passing in three arguments to the constructor. The first argument is the job name and refers to this job when using one of the `Scheduler` interface's administration methods, such as `pauseJob()`, which allows a particular job to be paused. The second argument is the group name, for which we are using the default. Group names can be used to refer to a group of jobs together, perhaps to pause them all using `Scheduler.pauseJobGroup()`. You should note that job names are unique within a group. The third and final argument is the `Class` that implements this particular job.

With the `JobDetail` instance created, we now move on to create a `Trigger`. In this example, we use the `SimpleTrigger` class, which provides `JDK Timer`-style trigger behavior. The first and second arguments passed to the `SimpleTrigger` constructor are the trigger name and group name, respectively. Both of these arguments perform similar functions for a `Trigger` as they do for a `JobDetail`. Also note that trigger names must be unique within a group, as otherwise, an exception is raised. The third and fourth arguments, both of type `Date`, are the start and end dates for this `Trigger`. By specifying `null` for the end date, we are saying there is no end date. The ability to specify an end date for a trigger is not available when you are using `Timer`. The next argument is the repeat count, which allows you to specify the maximum number of times the `Trigger` can fire. We use the constant `REPEAT_INDEFINITELY` to allow the `Trigger` to fire without a limit. The final argument is the interval between `Trigger` firings and is defined in milliseconds. We have defined an interval of 3 seconds.

The final step in this example is to schedule the job with a call to `Scheduler.schedule()` that passes in the `JobDetail` instance and the `Trigger`. If you run this application, you will see the familiar stream of “Hello, World” messages appearing gradually in your console.

Using JobDataMaps

In the previous example, all information for the job execution was contained in the job itself. However, you can pass state into the job using the `JobDetail` or `Trigger` class. Each instance of `JobDetail` and `Trigger` has an associated `JobDataMap` instance, which implements `Map`, and allows you to pass in job data in key/value pairs. Your jobs can modify data in the `JobDataMap` to allow for the passing of data between subsequent executions of the job. However, there are some considerations related to job persistence when using this approach. We discuss these later in the “About Job Persistence” section.

Storing data about a `Trigger` is useful when you have the same job scheduled with multiple `Trigger` implementations and want to provide the job with different data on each independent triggering. Entries of this map are made available via the `JobDataMap` on the `JobExecutionContext` that can be retrieved via the `getMergedJobDataMap()` method. As the method name suggests, the `JobDataMap` on the `JobExecutionContext` is a merge of the `JobDataMap` found on the `JobDetail` and the `JobDataMap` found on the `Trigger`, whereas data stored in the `Trigger` overrides data stored on the `JobDetail`.

In Listing 12-16, you can see an example of a `Job` that uses data contained in the merged `JobDataMap` to perform its processing.

Listing 12-16. Using the JobDataMap

```
package com.apress.prospring2.ch12.quartz;

import java.util.Map;

import org.quartz.Job;
import org.quartz.JobExecutionContext;
import org.quartz.JobExecutionException;

public class MessageJob implements Job {

    public void execute(JobExecutionContext context) throws JobExecutionException {
```

```

        Map properties = context.getMergedJobDataMap();

        System.out.println("Previous Fire Time: "
            + context.getPreviousFireTime());
        System.out.println("Current Fire Time: " + context.getFireTime());
        System.out.println("Next Fire Time: " + context.getNextFireTime());
        System.out.println(properties.get("message"));
        System.out.println(properties.get("jobDetailMessage"));
        System.out.println(properties.get("triggerMessage"));
        System.out.println("");
    }
}

```

From the merged `JobDataMap`, we are able to extract the Objects that are keyed as `message`, `jobDetailMessage`, and `triggerMessage` and write them to `stdout`. Also notice that we are able to get information about the previous, current, and next execution of this job from the `JobExecutionContext`.

In Listing 12-17, you can see an example of how you populate the `JobDataMap` on `JobDetail` with data when scheduling the Job.

Listing 12-17. *Adding Data to the JobDetail JobDataMap*

```

package com.apress.prospring2.ch12.quartz;

import org.quartz.Scheduler;
import org.quartz.SimpleTrigger;
import org.quartz.JobDetail;
import org.quartz.Trigger;
import org.quartz.impl.StdSchedulerFactory;

import java.util.Map;
import java.util.Date;

public class JobDetailMessageScheduling {

    public static void main(String[] args) throws Exception {
        Scheduler scheduler = new StdSchedulerFactory().getScheduler();
        scheduler.start();

        JobDetail jobDetail = new JobDetail("messageJob",
            Scheduler.DEFAULT_GROUP, MessageJob.class);

        Map map = jobDetail.getJobDataMap();
        map.put("message", "This is a message from Quartz");
        map.put("jobDetailMessage", "A jobDetail message");

        Trigger trigger = new SimpleTrigger("simpleTrigger",
            Scheduler.DEFAULT_GROUP, new Date(), null,
            SimpleTrigger.REPEAT_INDEFINITELY, 3000);

        scheduler.scheduleJob(jobDetail, trigger);
    }
}

```

You will recognize much of this code from the example in Listing 12-15, but notice that once the `JobDetail` instance has been created here, we access the `JobDataMap` and add two messages to it, keyed as `message` and `jobDetailMessage`. If you run this example and leave it running for a few iterations, you end up with output similar to this:

```

Previous Fire Time: null
Current Fire Time: Tue Oct 23 11:02:19 BST 2007
Next Fire Time: Tue Oct 23 11:02:22 BST 2007
This is a message from Quartz
A jobDetail message
null

```

```

Previous Fire Time: Tue Oct 23 11:02:19 BST 2007
Current Fire Time: Tue Oct 23 11:02:22 BST 2007
Next Fire Time: Tue Oct 23 11:02:25 BST 2007
This is a message from Quartz
A jobDetail message
null

```

```

Previous Fire Time: Tue Oct 23 11:02:22 BST 2007
Current Fire Time: Tue Oct 23 11:02:25 BST 2007
Next Fire Time: Tue Oct 23 11:02:28 BST 2007
This is a message from Quartz
A jobDetail message
null

```

You can see that both messages contained in the `JobDataMap` are written to stdout after the information about the execution times of the previous, current, and next execution is displayed.

Listing 12-18 gives you an example of also providing data on the `Trigger` and the effect that merging the two `JobDataMap` instances has on the values.

Listing 12-18. *Using the `JobDataMap` on the `Trigger`*

```

package com.apress.prospring2.ch12.quartz;

import org.quartz.JobDetail;
import org.quartz.Scheduler;
import org.quartz.SimpleTrigger;
import org.quartz.Trigger;
import org.quartz.impl.StdSchedulerFactory;

import java.util.Date;

public class TriggerMessageScheduling {

    public static void main(String[] args) throws Exception {
        Scheduler scheduler = new StdSchedulerFactory().getScheduler();
        scheduler.start();

        JobDetail jobDetail = new JobDetail("triggerMessageJob",
            Scheduler.DEFAULT_GROUP, MessageJob.class);
        jobDetail.getJobDataMap().put("message", "This is a message from Quartz");
        jobDetail.getJobDataMap().put("jobDetailMessage", "My job details data.");

        Trigger trigger = new SimpleTrigger("simpleTrigger",
            Scheduler.DEFAULT_GROUP, new Date(), null,
            SimpleTrigger.REPEAT_INDEFINITELY, 3000);
        trigger.getJobDataMap().put("message", "Message from Trigger");
        trigger.getJobDataMap().put("triggerMessage", "Another trigger message.");
    }
}

```

```

        scheduler.scheduleJob(jobDetail, trigger);
    }
}

```

As you can see, the `JobDetail` is configured exactly as before. We simply add two messages to the `Trigger`: a keyed message and a trigger message. Running this example produces output similar to the following:

```

Previous Fire Time: null
Current Fire Time: Tue Oct 23 11:14:22 BST 2007
Next Fire Time: Tue Oct 23 11:14:25 BST 2007
Message from Trigger
My job details data.
Another trigger message.

Previous Fire Time: Tue Oct 23 11:14:22 BST 2007
Current Fire Time: Tue Oct 23 11:14:25 BST 2007
Next Fire Time: Tue Oct 23 11:14:28 BST 2007
Message from Trigger
My job details data.
Another trigger message.

Previous Fire Time: Tue Oct 23 11:14:25 BST 2007
Current Fire Time: Tue Oct 23 11:14:28 BST 2007
Next Fire Time: Tue Oct 23 11:14:31 BST 2007
Message from Trigger
My job details data.
Another trigger message.

```

Note that the value of the key message is what we added to the `Trigger JobDataMap`, not what we defined in the `JobDetail`.

As you will see shortly, when using Spring to configure Quartz scheduling, you can create the `JobDataMap` in your Spring configuration file, allowing you to externalize all Job configuration completely.

Using the CronTrigger

In the previous examples, we used the `SimpleTrigger` class, which provides trigger functionality very similar to that of the `JDK Timer` class. However, Quartz excels in its support for complex trigger expressions using the `CronTrigger`. `CronTrigger` is based on the Unix cron daemon, a scheduling application that supports a simple, yet extremely powerful, trigger syntax. Using `CronTrigger`, you can quickly and accurately define trigger expressions that would be extremely difficult or impossible to do with the `SimpleTrigger` class. For instance, you can create a trigger that says, “Fire every 5 seconds of every minute, starting at the third second of the minute, but only between the hours of 2:00 and 5:00 p.m.” Or it could say, “Fire on the last Friday of every month.”

A `CronTrigger` syntax expression, referred to as a cron expression, contains six required components and one optional component. A cron expression is written on a single line, and each component is separated from the next by a space. Only the last, or rightmost, component is optional. Table 12-1 describes the cron components in detail.

Table 12-1. *Components of a cron Expression*

Position	Meaning	Allowed Special Characters
1	Seconds (0–59)	, , - , * , and /
2	Minutes (0–59)	, , - , * , and /
3	Hours (0–23)	, , - , * , and /
4	Day of the month (1–31)	, , - , * , / , ? , L , W , and C
5	Month (either JAN–DEC or 1–12)	, , - , * , and /
6	Day of the week (either SUN–SAT or 1–7)	, , - , * , / , ? , L , C , and #
7	Year (optional, 1970–2099), when empty, full range is assumed	, , - , * , and /

Each component accepts the typical range of values that you would expect, such as 0–59 for seconds and minutes and 1–31 for the day of the month. For the month and day of the week components, you can use numbers, such as 1–7 for day of the week, or text such as SUN–SAT.

Each field also accepts a given set of special symbols, so placing an asterisk (*) in the hours component means “every hour,” and using an expression such as 6L in the day of the week component means “last Friday of the month.” Table 12-2 describes cron wildcards and special characters in detail.

Table 12-2. *cron Expression Wildcards and Special Characters*

Special Character	Description
*	Any value. This special character can be used in any field to indicate that the value should not be checked. Therefore, our example cron expression will be fired on every day of the month, every month, and every day of the week between 1970 and 2099.
?	No specific value. This special character is usually used with other specific values to indicate that a value must be present but will not be checked.
-	Range. For example, 10–12 in the hours field means hours 10:00, 11:00, and 12:00 a.m.
,	List separator. Allows you to specify a list of values, such as MON, TUE, WED in the day of the week field.
/	Increments. This character specifies increments of a value. For example, 0/1 in the minute field means the job should run on every 1-minute increment of the minute field, starting from 0.
L	An abbreviation for “last.” The meaning of L is a bit different in the day of the month field than for the day of the week. When used in the day of the month field, it means the last day of the month (March 31, February 28 or 29, and so on). When used in the day of the week field, it has the same value as 7—Saturday. The L special character is most useful when you use it with a specific day of the week value. For example, 6L in the day of the week field means the last Friday of each month.
W	W is only allowed for the day of the month field and specifies the nearest weekday (Monday–Friday) to the given day of the same month. Set the value to 7W, and the trigger will be fired on the sixth day if the seventh happens to be a Saturday. If the seventh day is a Sunday, the trigger fires on Monday, the eighth day. Note that a trigger due on Saturday the first will actually fire on the third.

Continued

Table 12-2. *Continued*

Special Character	Description
#	This value is allowed only for the day of the week field, and it specifies the <i>n</i> th day in a month. For example 1#2 means the first Monday of each month.
C	The calendar value. This is allowed for the day of the month and day of the week fields. The values of days are calculated against a specified calendar. Specifying 20C in the day of the month field fires the trigger on the first day included in the calendar on or after the twentieth day. Specifying 6C in the day of the week field is interpreted as the first day included in the calendar on or after Friday.

The last thing to bear in mind when writing cron expressions is daylight saving time. Changes because of daylight saving time may cause a trigger to fire twice in spring or to never fire in autumn.

There are many more permutations for cron expressions than we can discuss here; you can find a detailed description of cron syntax in the Javadoc for the `CronTrigger` class.

Listing 12-19 shows an example of the `CronTrigger` class in action.

Listing 12-19. *Using the CronTrigger Class*

```
package com.apress.prospring2.ch12.quartz;

import java.util.Map;

import org.quartz.CronTrigger;
import org.quartz.JobDetail;
import org.quartz.Scheduler;
import org.quartz.Trigger;
import org.quartz.impl.StdSchedulerFactory;

public class CronTriggerExample {

    public static void main(String[] args) throws Exception {
        Scheduler scheduler = new StdSchedulerFactory().getScheduler();
        scheduler.start();

        JobDetail jobDetail = new JobDetail("messageJob",
            Scheduler.DEFAULT_GROUP, MessageJob.class);

        Map map = jobDetail.getJobDataMap();
        map.put("message", "This is a message from Quartz");

        String cronExpression = "3/5 * 14,15,16,17 * * ?";

        Trigger trigger = new CronTrigger("cronTrigger",
            Scheduler.DEFAULT_GROUP, cronExpression);

        scheduler.scheduleJob(jobDetail, trigger);
    }
}
```

Much of this code should look familiar to you; the only major difference here is that we use the cron expression. The actual creation of the `CronTrigger` class is very similar to the creation of the `SimpleTrigger` class in that you have a name and a group name. To help you understand the cron expression in the example, we'll break it down into components.

The first component, 3/5, means every 5 seconds starting at the third second of the minute. The second component, *, simply says every minute. The third component, 14, 15, 16, 17, restricts this trigger to running between 2:00 and 5:59 p.m.—that is, when the time begins with 14, 15, 16, or 17. The next two components are both wildcards saying that this trigger can run in any month or any year. The final component uses the question mark wildcard, ?, to indicate that this trigger can run on any day of the week. This expression has the net effect of firing every 5 seconds, starting on the third second of the minute, but only between 2:00 and 5:59 p.m.

If you run this example, depending on the time of day, you see either a blank screen or the ever increasing list of “Hello, World” printouts. Try modifying the first component in the expression to change the frequency of the trigger or at which second in the minute the trigger starts. You should also try modifying other components to see what effects you get.

The CronTrigger class is great for almost all trigger requirements. However, expressions can quickly become convoluted when you need to program exceptions to the rule. For instance, consider a process that checks a task list for a user every Monday, Wednesday, and Friday at 11:00 a.m. and 3:00 p.m. Now, consider what happens when you want to prevent this trigger from firing when the user is on vacation. Thankfully, Quartz provides support for this via the Calendar interface. Using the Calendar interface, you can accurately define a period that should either be explicitly included or explicitly excluded from a trigger’s normal schedule. Quartz comes with six implementations of Calendar, one of which is the HolidayCalendar that stores a list of days to be excluded from a trigger’s schedule. Listing 12-20 shows a modification of the previous example that uses a HolidayCalendar to exclude December 25, 2007.

Listing 12-20. *Explicitly Excluding Dates with HolidayCalendar*

```
package com.apress.prospring2.ch12.quartz;

import java.util.Calendar;
import java.util.Map;

import org.quartz.CronTrigger;
import org.quartz.JobDetail;
import org.quartz.Scheduler;
import org.quartz.Trigger;
import org.quartz.impl.StdSchedulerFactory;
import org.quartz.impl.calendar.HolidayCalendar;

public class CronWithCalendarExample {

    public static void main(String[] args) throws Exception {
        Scheduler scheduler = new StdSchedulerFactory().getScheduler();
        scheduler.start();

        // create a calendar to exclude a particular date
        Calendar cal = Calendar.getInstance();
        cal.set(2007, Calendar.DECEMBER, 25);

        HolidayCalendar calendar = new HolidayCalendar();
        calendar.addExcludedDate(cal.getTime());

        // add to scheduler
        scheduler.addCalendar("xmasCalendar", calendar, true, false);

        JobDetail jobDetail = new JobDetail("messageJob",
            Scheduler.DEFAULT_GROUP, MessageJob.class);
```

```

    Map map = jobDetail.getJobDataMap();
    map.put("message", "This is a message from Quartz");

    String cronExpression = "3/5 * 14,15,16,17 * * ?";

    Trigger trigger = new CronTrigger("cronTrigger",
        Scheduler.DEFAULT_GROUP, cronExpression);

    trigger.setCalendarName("xmasCalendar");

    scheduler.scheduleJob(jobDetail, trigger);
}
}

```

Here, you can see that we create an instance of `HolidayCalendar`, and using the `addExcludedDate()` method, we exclude December 25. With the `Calendar` instance created, we add the `Calendar` to the `Scheduler` using the `addCalendar()` method, giving it a name of `xmasCalendar`. Later, before adding the `CronTrigger`, we associate it with `xmasCalendar`. Using this approach saves you from having to create complex cron expressions just to exclude a few arbitrary dates.

About Job Persistence

Quartz provides support for Job persistence, allowing you to add jobs at runtime or make changes to existing jobs and persist these changes and additions for subsequent executions of the Job. Central to this concept is the `JobStore` interface, implementations of which are used by Quartz when it is performing persistence. By default, Quartz uses the `RAMJobStore` implementation, which simply stores Job instances in memory. Other available implementations are `JobStoreCMT` and `JobStoreTX`. Both of these classes persist job details using a configured `DataSource` and support the creation and modification of jobs as part of a transaction. The `JobStoreCMT` implementation is intended to be used in an application server environment and takes part in container-managed transactions. For stand-alone applications, you should use the `JobStoreTX` implementation. Spring provides its own `LocalDataSourceJobStore` implementation of `JobStore`, which can take part in Spring-managed transactions. We will take a look at this implementation when we discuss Spring support for Quartz.

Earlier on, you saw how you can modify the contents of the `JobDataMap` to pass information between different executions of the same Job. However, if you try to run that example using a `JobStore` implementation other than `RAMJobStore`, you will be surprised to see that it doesn't work. The reason for this is that Quartz supports the notion of stateless and stateful jobs. When using the `RAMJobStore` and modifying the `JobDataMap`, you are actually modifying the store directly, so the type of Job is unimportant, but this is not the case when you are using implementations other than `RAMJobStore`. A stateless Job only has the data in the `JobDataMap` persisted when it is added to the `Scheduler`, whereas stateful Jobs have their `JobDataMap` persisted after every execution. To mark a Job as stateful, implement the `StatefulJob` interface instead of the `Job` interface. `StatefulJob` is a subinterface of `Job`, so you do not need to implement `Job` as well. You should also be aware that any data you place in the `JobDataMap` when using Job persistence must be serializable, because Quartz writes the `JobDataMap` as a serialized blob to the database.

Quartz Support in Spring

Spring's Quartz integration follows a similar pattern to the integration with `Timer` in that it allows you to configure your job scheduling fully within the Spring configuration file. In addition to this, Spring provides further classes to integrate with the Quartz `JobStore`, thus allowing you to configure Job persistence in your configuration and for Job modification to take part in Spring-managed transactions.

Scheduling a Job with Spring

As you would expect, much of the code you need to schedule a Quartz Job using Spring goes into the Spring configuration file. Indeed, you only need to load the `ApplicationContext` in your application for the configuration to take effect and for Spring to start the Scheduler automatically.

In Listing 12-21, you can see the configuration code required to configure the `MessageJob` class you saw in Listing 12-16 to run once every 3 seconds.

Listing 12-21. *Configuring Scheduling Declaratively*

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans-2.0.xsd">

  <bean id="job" class="org.springframework.scheduling.quartz.JobDetailBean">
    <property name="jobClass"
      value="com.apress.prospring2.ch12.quartz.MessageJob"/>
    <property name="jobDataAsMap">
      <map>
        <entry key="message"
          value="This is a message from the Spring config file!"/>
      </map>
    </property>
  </bean>

  <bean id="trigger"
    class="org.springframework.scheduling.quartz.SimpleTriggerBean">
    <property name="jobDetail" ref="job"/>
    <property name="startDelay" value="1000"/>
    <property name="repeatInterval" value="3000"/>
    <property name="jobDataAsMap">
      <map>
        <entry key="triggerMessage"
          value="Trigger message from the Spring config file!"/>
      </map>
    </property>
  </bean>

  <bean id="schedulerFactory"
    class="org.springframework.scheduling.quartz.SchedulerFactoryBean">
    <property name="triggers">
      <list>
        <ref local="trigger"/>
      </list>
    </property>
  </bean>
</beans>
```

Here, you can see that we use the `JobDetailBean` class, which extends the `JobDetail` class, to configure the job data in a declarative manner. The `JobDetailBean` provides more JavaBean-style properties that are accessible by Spring, and it also provides sensible defaults for properties that you usually have to specify yourself. For instance, notice that we did not specify a job name or a group name. By default, the `JobDetailBean` uses the ID of the `<bean>` tag as the job name and the default group of the Scheduler as the group name. Notice that we are able to add data to the `JobDataMap` property using the `jobDataAsMap` property. The name of this property is not a typographical error—

you can't add directly to the `jobDataMap` property. It is of type `JobDataMap`, and this type is not supported in Spring configuration files.

With the `JobDetailBean` configured, the next step is to create a trigger. Spring offers two classes, `SimpleTriggerBean` and `CronTriggerBean`, that wrap the `SimpleTrigger` and `CronTrigger` classes, allowing you to configure them declaratively and to associate them with a `JobDetailBean`—all within your configuration file. Notice that in Listing 12-21, we defined a starting delay of 1 second and a repeat interval of 3 seconds. By default, the `SimpleTriggerBean` sets the repeat count to infinity.

The final piece of configuration you need is the `SchedulerFactoryBean`. By default, the `SchedulerFactoryBean` creates an instance of `StdSchedulerFactory`, which, in turn, creates the Scheduler implementation. You can override this behavior by setting the `schedulerFactoryClass` property to the name of a class that implements `SchedulerFactory`, which you wish to use in place of `StdSchedulerFactory`. The only property that you need to configure scheduling is the `triggers` property, which accepts a List of `TriggerBean` elements.

Because all of the job scheduling configuration is contained in the configuration, you need very little code to actually start the Scheduler and execute the Job instances. In fact, all you need to do is create the `ApplicationContext`, as shown in Listing 12-22.

Listing 12-22. *Testing Declarative Quartz Configuration*

```
package com.apress.prospring2.ch12.quartz.spring;

import org.springframework.context.ApplicationContext;
import org.springframework.context.support.FileSystemXmlApplicationContext;

public class SimpleSpringQuartzIntegrationExample {

    public static void main(String[] args) {
        ApplicationContext ctx = new FileSystemXmlApplicationContext(
            "./ch12/src/conf/quartz-simple.xml ");
    }
}
```

As you can see, this class does nothing more than create an instance of `ApplicationContext` using the configuration shown in Listing 12-21. If you run this application and leave it running for a few iterations, you end up with something like this:

```
Previous Fire Time: null
Current Fire Time: Tue Oct 23 11:24:31 BST 2007
Next Fire Time: Tue Oct 23 11:24:34 BST 2007
This is a message from the Spring configuration file!
null
Trigger message from the Spring configuration file!

Previous Fire Time: Tue Oct 23 11:24:31 BST 2007
Current Fire Time: Tue Oct 23 11:24:34 BST 2007
Next Fire Time: Tue Oct 23 11:24:37 BST 2007
This is a message from the Spring configuration file!
null
Trigger message from the Spring configuration file!

Previous Fire Time: Tue Oct 23 11:24:34 BST 2007
Current Fire Time: Tue Oct 23 11:24:37 BST 2007
Next Fire Time: Tue Oct 23 11:24:40 BST 2007
This is a message from the Spring configuration file!
null
Trigger message from the Spring configuration file!
```

Notice that it is running just like it was for the previous `MessageJob` example, but the messages displayed are the messages configured in the Spring configuration file.

Using Persistent Jobs

One of the great features of Quartz is its ability to create stateful, persistent jobs. This opens up some great functionality that is not available when you are using `Timer`-based scheduling. With persistent jobs, when you can add jobs to Quartz at runtime, they will still be in your application after a restart. Plus, you can modify the `JobDataMap` passed between executions of a `Job`, and the changes will still be in effect after a restart.

In this example, we are going to schedule two jobs, one using Spring configuration mechanisms and one at runtime. We'll see how the Quartz persistence mechanism copes with changes to the `JobDataMap` for these jobs and what happens in subsequent executions of the application.

To start with, you need to create a database in which Quartz can store the `Job` information. In the Quartz distribution—we used version 1.6.0—you will find a selection of database scripts for a variety of different RDBMS flavors. For the example here, we use Oracle, but you should not encounter problems using a different database as long as Quartz has a database script for it. For version 1.6.0, try the `docs/dbTables` subfolder of the Quartz distribution. Once you have located the script for your database, execute it against your database and verify that 12 tables, each with the prefix `qrtz`, have been created.

Next, create your test `Job`. Because we want to make changes to `JobDataMap` during `Job` execution, we need to flag to Quartz that it should treat this as a stateful `Job`. We do this by implementing the `StatefulJob` interface rather than the `Job` interface. This is shown in Listing 12-23.

Listing 12-23. *Creating a Stateful Job*

```
package com.apress.prospring2.ch12.quartz.spring;

import java.util.Map;

import org.quartz.JobExecutionContext;
import org.quartz.JobExecutionException;
import org.quartz.StatefulJob;

public class PersistentJob implements StatefulJob {

    public void execute(JobExecutionContext context)
        throws JobExecutionException {
        Map map = context.getJobDetail().getJobDataMap();
        System.out.println "[" + context.getJobDetail().getName() + "]"
            + map.get("message");
        map.put("message", "Updated Message");
    }
}
```

The `StatefulJob` interface does not declare additional methods for your class to implement; it is simply a marker telling Quartz that it should persist the `JobDetail` after every execution. Here, you can see that we display the message that is stored in the `JobDataMap` along with the name of the `Job`.

The next steps are to configure the `Job` in Spring and configure the `Scheduler` with a `DataSource` it can use for persistence, as shown in Listing 12-24.

Listing 12-24. *Configuring Quartz Persistence in Spring*

```

<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
       xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
       xsi:schemaLocation="http://www.springframework.org/schema/beans
                           http://www.springframework.org/schema/beans/spring-beans-2.0.xsd">

    <bean id="job" class="org.springframework.scheduling.quartz.JobDetailBean">
        <property name="jobClass"
            value="com.apress.prospring2.ch12.quartz.spring.PersistentJob"/>
        <property name="jobDataAsMap">
            <map>
                <entry key="message" value="Original Message"/>
            </map>
        </property>
    </bean>

    <bean id="dataSource"
        class="org.springframework.jdbc.datasource.SingleConnectionDataSource">
        <property name="driverClassName" value="oracle.jdbc.driver.OracleDriver"/>
        <property name="url"
            value="jdbc:oracle:thin:@oracle.devcake.co.uk:1521:INTL"/>
        <property name="username" value="PROSPRING"/>
        <property name="password" value="x*****6"/>
    </bean>

    <bean id="trigger"
        class="org.springframework.scheduling.quartz.SimpleTriggerBean">
        <property name="jobDetail" ref="job"/>
        <property name="startDelay" value="1000"/>
        <property name="repeatInterval" value="3000"/>
    </bean>

    <bean id="schedulerFactory"
        class="org.springframework.scheduling.quartz.SchedulerFactoryBean">
        <property name="triggers">
            <list>
                <ref local="trigger"/>
            </list>
        </property>
        <property name="dataSource" ref="dataSource"/>
    </bean>
</beans>

```

You will recognize much of this configuration code from Listing 12-21; the important part here is the `dataSource` bean. In this code, we use the Spring class `SingleConnectionDataSource`; this `DataSource` implementation is handy for testing, but never use it in production (check the Javadoc for this class if you are unsure why). Also, remember, you need to modify the connection details in the configuration as appropriate for your environment. For more details on configuring other `DataSources` with Spring, see Chapter 8.

Using the configured `dataSource` bean, we set the `dataSource` property of the `SchedulerFactoryBean`. By doing this, we instruct Spring to create a `Scheduler` that is configured to persist Job data using the given `DataSource`. Internally, this is achieved using Spring's own `JobStore` implementation, `LocalDataSourceJobStore`.

With the configuration complete, all that remains is to load it in an application and add another Job to the Scheduler at runtime. Listing 12-25 shows the code for this.

Listing 12-25. *Testing Job Persistence*

```
package com.apress.prospring2.ch12.quartz.spring;

import java.util.Date;

import org.quartz.JobDetail;
import org.quartz.Scheduler;
import org.quartz.SimpleTrigger;
import org.quartz.Trigger;
import org.springframework.context.ApplicationContext;
import org.springframework.context.support.FileSystemXmlApplicationContext;

public class SpringWithJobPersistence {

    public static void main(String[] args) throws Exception {
        ApplicationContext ctx = new FileSystemXmlApplicationContext(
            "./ch12/src/conf/quartzPersistent.xml");

        // get the scheduler
        Scheduler scheduler = (Scheduler) ctx.getBean("schedulerFactory");

        JobDetail job = scheduler.getJobDetail("otherJob",
            Scheduler.DEFAULT_GROUP);

        if (job == null) {
            // the job has not yet been created
            job = (JobDetail) ctx.getBean("job");
            job.setName("otherJob");
            job.getJobDataMap().put("message", "This is another message");

            Trigger trigger = new SimpleTrigger("simpleTrigger",
                Scheduler.DEFAULT_GROUP, new Date(), null,
                SimpleTrigger.REPEAT_INDEFINITELY, 3000);

            scheduler.scheduleJob(job, trigger);
        }
    }
}
```

This code requires little explanation. However, note that before we schedule the second job, we check to see if it already exists using the `Scheduler.getJobDetail()` method. This prevents us from overwriting the Job on subsequent runs of the application.

The first time you run this example, you get output something like this:

```
[otherJob]This is another message
[job]Original Message
[otherJob]Updated Message
[job]Updated Message
[otherJob]Updated Message
[job]Updated Message
```

As you can see, the first time each Job executes, the message displayed is the original message configured in the JobDataMap when the Job was scheduled. On subsequent executions, each Job displays the updated message that was set during the previous execution. If you stop the application and restart it, you see something slightly different:

```
[otherJob]Updated Message
[job]Updated Message
[otherJob]Updated Message
[job]Updated Message
[otherJob]Updated Message
[job]Updated Message
```

This time, you can see that, because the Job data was persisted, you do not need to re-create the second Job, and the JobDataMap accurately reflects changes that were made during the last run of the application.

Scheduling Arbitrary Jobs with Quartz

Like the Timer-based scheduling classes, Spring provides the ability to schedule the execution of arbitrary methods using Quartz. We won't go into detail on this, because it works in an almost identical manner to the Timer approach. Instead of using `MethodInvokingTimerTaskFactoryBean`, you use `MethodInvokingJobDetailFactoryBean`, and instead of automatically creating `TimerTask` implementations, you automatically create `JobDetail` ones.

Job Scheduling Considerations

If you are going to be adding job scheduling to your application, you should bear in mind a few considerations when you choose a scheduler and a scheduling approach.

Choosing a Scheduler

The first decision you have to make when adding scheduling to your application is which scheduler to use. This choice is actually quite easy. If you have only very simple scheduling requirements or are restricted in the external libraries that you can package with your application, you should use Timer-based scheduling. Otherwise, use Quartz.

Even if you find that your requirements are simple, you might want to go with Quartz, especially if you have to create an explicit Job implementation. If you use Quartz from the outset, if your requirements become more advanced, you can easily add persistence, transactions, or more complex triggers without having to change a `TimerTask` to a Job. In general, we have found that using Quartz for all of our scheduling allows us to become familiar with a single scheduling approach and saves our developers from having to worry about two different approaches when one provides everything they need.

Packaging Job Logic Separately from the Job Class

A common approach that we see many developers take when adding scheduling to an application is to place business logic inside a Job or `TimerTask`. Generally, this is a bad idea. In many cases, you need to have scheduled tasks available for execution on demand, which requires the logic to be separate from the scheduling framework.

Also, you should not unnecessarily couple your business logic to a particular scheduler. We have found that a better approach is to keep business logic in separate classes and either create

a simple wrapper around those classes that is specific to your scheduler or, preferably, use the appropriate `MethodInvoker*FactoryBean` to create the wrapper for you.

Task Execution and Thread Pooling

So far in this chapter, we have discussed various ways of scheduling jobs to be executed at a specific point in time, at defined intervals, or using a combination of both times and intervals. Now, we are going to look at another way to schedule jobs in Spring that depend less on a specific time or interval than on immediate or event-triggered execution.

For example, think of a web server handling incoming requests. A simple approach for building such a server application would be to process each job in a new thread. Depending on the server you are building and its environment, this might work absolutely fine. But as the creation of a thread needs time and system resources, you might end up spending more time creating and destroying threads than executing jobs, not to mention that you might run out of system resources. To run stably, a server needs some way of managing how much can be done at the same time. The concept of thread pools and work queue offers just this.

The `java.util.concurrent` Package

One welcome addition to Java 5 was the `java.util.concurrent` package based on Doug Lea's `util.concurrent` package, a library offering efficient and well-tested tools to simplify the development of multithreaded applications. This package provides the `Executor` interface, which defines only one method `execute(Runnable command)` to execute `Runnable` tasks. It abstracts the submission of tasks away from the details of how they are run. Implementations of this interface offer all sorts of execution policy: thread-per-task, thread pooling, or even synchronous execution just to name a few (you can find some implementations of these in the Javadoc for the `Executor` interface).

To give you a little example of how this interface and its subinterface `ExecutorService` can be used, we will first create a task to execute using a slightly amended version of our former `HelloWorldTask` from Listing 12-1. We could just as well use the `HelloWorldTask` straight away, as it extends `TimerTask` that implements `Runnable`, but we wouldn't be able to see the task scheduling differences between various `Executor` implementations.

Listing 12-26. *HelloWorldCountDownTask*

```
package com.apress.prospring2.ch12;

public class HelloWorldCountDownTask implements Runnable {

    private String name;
    private int count = 4;

    public HelloWorldCountDownTask(String name) {
        this.name = name;
    }

    public void run() {
        while (count > 0) {
            count--;
            if (count == 0) {
                System.out.println(name + " says 'Hello World!');
            } else {
                System.out.println(name + ": " + count);
                Thread.yield();
            }
        }
    }
}
```

```

    }
}

```

All this task does is print out a countdown from 3, calling the `Thread.yield()` method afterward to pause this thread's executions and allow other threads to be executed. As a last statement, the task will say hello to the world and finish execution.

Next, as shown in Listing 12-27, we are going to use an `ExecutorService` implementation to schedule and execute this task.

Listing 12-27. *Use `ExecutorService` to Schedule and Execute a Task*

```

package com.apress.prospring2.ch12.executor;
import com.apress.prospring2.ch12.HelloWorldCountDownTask;
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;

public class ExecutorServiceExample {

    public static void main(String[] args) {
        ExecutorService service = Executors.newFixedThreadPool(2);

        service.execute(new HelloWorldCountDownTask("Anna"));
        service.execute(new HelloWorldCountDownTask("Beth"));
        service.execute(new HelloWorldCountDownTask("Charlie"));
        service.execute(new HelloWorldCountDownTask("Daniel"));

        service.shutdown();
    }
}

```

The `java.util.concurrent.Executors` class provides convenient factory and utility methods for `Executor` and `ExecutorService` classes. We are using its `newFixedThreadPool()` method to retrieve a `ThreadPoolExecutor` with a fixed number of two threads in the pool. We then submit four tasks to execution and call the `ExecutorService`'s `shutdown()` method to shut down the `ExecutorService` after all tasks have been executed. After calling this method, no further tasks can be added to the service. Running this example will print out the following:

```

Anna: 3
Anna: 2
Beth: 3
Anna: 1
Anna says 'Hello World!
Charlie: 3
Beth: 2
Charlie: 2
Beth: 1
Charlie: 1
Beth says 'Hello World!
Daniel: 3
Charlie says 'Hello World!
Daniel: 2
Daniel: 1
Daniel says 'Hello World!

```

Note that there are only two tasks being executed at the same time, and task Charlie is only getting executed after task Anna has finished. Try a different number of threads in the pool or a different Executor implementation, and you will find the printout to be different.

Spring's TaskExecutor Abstraction

Since version 2.0, Spring has offered an abstraction to the previously discussed Java 5 Executor framework. Identical to the `java.util.concurrent.Executor` interface, the `TaskExecutor` interface defines only the single method `execute(Runnable command)`. Intended to be used internally in other Spring components such as asynchronous JMS and JCA environment support, it now lets you add thread pooling behavior to your own application without creating the need for Java 5.

Spring comes with a variety of `TaskExecutor` implementations, which are described in Table 12-3.

Table 12-3. *Spring's TaskExecutor Implementations*

Implementation	Description
<code>SimpleAsyncTaskExecutor</code>	This implementation provides asynchronous threading with a new thread per invocation policy. It also allows setting a concurrency limit that will block further invocations.
<code>SyncTaskExecutor</code>	When you choose this implementation task, execution happens synchronously in the calling thread.
<code>ConcurrentTaskExecutor</code>	This class implements Spring's <code>SchedulingTaskExecutor</code> interface as well Java 5's <code>java.util.concurrent.Executor</code> interface and acts as a wrapper for the latter.
<code>SimpleThreadPoolTaskExecutor</code>	This is a subclass of Quartz's <code>SimpleThreadPool</code> and useful if a thread pool needs to be shared by Quartz and non-Quartz components.
<code>ThreadPoolTaskExecutor</code>	Behaving similarly to the <code>ConcurrentTaskExecutor</code> , it exposes the <code>java.util.concurrent.ThreadPoolExecutor</code> parameters as bean properties (it needs Java 5).
<code>TimerTaskExecutor</code>	This implementation uses a <code>TimerTask</code> behind the scenes. Invocations are executed in a separate thread but synchronously within that thread.
<code>WorkManagerTaskExecutor</code>	This uses the CommonJ <code>WorkManager</code> implementation and implements the <code>WorkManager</code> interface.

The differences in how the various implementations work is easiest to see in a small example. Listing 12-28 shows the configuration of three `TaskExecutor` implementations in Spring: `ThreadPoolTaskExecutor`, `SyncTaskExecutor`, and `TimerTaskExecutor`.

Listing 12-28. *task-executor-context.xml*

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans-2.5.xsd">

  <bean id="threadPoolTaskExecutor"
    class="org.springframework.scheduling.concurrent.ThreadPoolTaskExecutor">
    <property name="corePoolSize" value="5"/>
```

```

        <property name="maxPoolSize" value="10"/>
        <property name="queueCapacity" value="25"/>
    </bean>

    <bean id="synchTaskExecutor"
        class="org.springframework.core.task.SyncTaskExecutor"/>

    <bean id="timerTaskExecutor"
        class="org.springframework.scheduling.timer.TimerTaskExecutor">
        <property name="delay" value="3000"/>
        <property name="timer" ref="timer"/>
    </bean>

    <bean id="timer" class="java.util.Timer"/>
</beans>

```

We can then use the defined beans to load them from the `ApplicationContext` and use them in a `TaskExecutorExample`, as shown in Listing 12-29.

Listing 12-29. *TaskExecutorExample*

```

package com.apress.prospring2.ch12.taskexecutor;

import org.springframework.core.task.TaskExecutor;
import com.apress.prospring2.ch12.HelloWorldCountDownTask;

public class TaskExecutorExample {

    private TaskExecutor taskExecutor;

    public TaskExecutorExample(TaskExecutor taskExecutor) {
        this.taskExecutor = taskExecutor;
    }

    public void executeTasks() {
        this.taskExecutor.execute(new HelloWorldCountDownTask("Anna"));
        this.taskExecutor.execute(new HelloWorldCountDownTask("Beth"));
        this.taskExecutor.execute(new HelloWorldCountDownTask("Charlie"));
        this.taskExecutor.execute(new HelloWorldCountDownTask("Daniel"));
    }
}

```

As you see, we use the same four `HelloWorldCountDownTask` instances as before. The generated output will highlight the different execution strategies. As expected, the `SyncTaskExecutorExample`, shown in Listing 12-30, executes the tasks synchronously:

Listing 12-30. *SyncTaskExecutorExample*

```

package com.apress.prospring2.ch12.taskexecutor;

import org.springframework.context.ApplicationContext;
import org.springframework.context.support.ClassPathXmlApplicationContext;
import org.springframework.core.task.TaskExecutor;

import java.io.IOException;

public class SynchTaskExecutorExample {

```

```

public static void main(String[] args) throws IOException {
    ApplicationContext ctx = new ClassPathXmlApplicationContext(
        "com/apress/prospring2/ch12/taskexecutor/task-executor-context.xml");

    TaskExecutor taskExecutor =
        (TaskExecutor)ctx.getBean("synchTaskExecutor", TaskExecutor.class);
    TaskExecutorExample example = new TaskExecutorExample(taskExecutor);
    example.executeTasks();
}
}

```

Running this code will create output like what's shown here:

```

Anna: 3
Anna: 2
Anna: 1
Anna says 'Hello World!
Beth: 3
Beth: 2
Beth: 1
Beth says 'Hello World!
Charlie: 3
Charlie: 2
Charlie: 1
Charlie says 'Hello World!
Daniel: 3
Daniel: 2
Daniel: 1
Daniel says 'Hello World!

```

If your application is running on Java 5 or higher, you can configure this abstraction to delegate to any of Java 5's implementations. Spring's `ThreadPoolTaskExecutor` enables you to configure a JDK 1.5 `ThreadPoolExecutor` through bean properties and exposes it as a Spring `TaskExecutor`. Additionally, Spring provides the `ConcurrentTaskExecutor` as an adapter class for other Java 5 `Executor` implementations, making upgrading from Java 1.4 even easier.

The adapter class implements both the `TaskExecutor` and `Executor` interfaces. Because the primary interface is the `TaskExecutor`, exception handling follows its contract. For example, when a task cannot be accepted for execution, a `SpringTaskRejectedExecutionException` is thrown rather than a `java.util.concurrent.RejectedExecutionException`.

A further convenient feature of the `TaskExecutor` interface is its wrapping of exceptions in runtime exceptions. When a task fails with an exception, the situation is usually considered fatal. Without the need or possibility to recover, the exception can go unchecked—your code stays more portable, and you can quite easily switch between `TaskExecutor` implementations.

Summary

In this chapter, we showed you various mechanisms for scheduling jobs with Spring. We looked at the basic support offered when you use JDK `Timer` and the more sophisticated support offered through Quartz. You saw how the different trigger types are used, and in particular, we explored the `CronTrigger` in Quartz as a means of creating complex schedules that match real-world scenarios.

Job scheduling is an important part of enterprise applications, and Spring provides excellent support for adding scheduling to your own applications. In the next chapter, we are going to examine Spring's support for sending e-mail messages.

